

Multiple Secret Leaders

Authors here

July 17, 2023

1 Some Notation

1. We will have n parties
2. We will have k leaders elected
3. We will have a “bid” published by a user $i \in [n]$ be denoted as b_i
4. We will denote a commitment as comm_i
5. We will denote some generic CRHF as h
6. We will say Enc_{TFHE} and Dec_{TFHE} for TFHE encoding and decoding respectively

2 Preliminaries

2.1 Threshold FHE

[JRS17, BGG⁺18] defines threshold FHE encryption. For the sake of completeness, we will define it here.

Definition 2.1 (TFHE [JRS17]). Let $P = \{P_1, \dots, P_N\}$ be a set of N parties and \mathbb{S} be a class of access structures on P . A TFHE scheme for \mathbb{S} is a tuple of PPT algorithms

$$(\text{TFHE.Setup}, \text{TFHE.Encrypt}, \text{TFHE.Eval}, \text{TFHE.PartDec}, \text{TFHE.FinDec})$$

such that the following specifications are met

- $(pk, sk_1, \dots, sk_N) \leftarrow \text{TFHE.Setup}(1^\lambda, 1^s, \mathbb{A})$: Takes as input a security parameter λ , a depth bound on the circuit, and a structure $\mathbb{A} \in \mathbb{S}$. Outputs a public key pk and a secret key sk_i for each party P_i .
- $\text{ct} \leftarrow \text{TFHE.Encrypt}(pk, \mu)$: Takes as input a public key and a message $\mu \in \{0, 1\}$ and outputs a ciphertext ct .
- $\hat{\text{ct}} \leftarrow \text{TFHE.Eval}(C, \text{ct}_1, \dots, \text{ct}_k)$: Takes as input a circuit C of depth at most d and k ciphertexts $\text{ct}_1, \dots, \text{ct}_k$. Outputs a ciphertext $\hat{\text{ct}} = C(\text{ct}_1, \dots, \text{ct}_k)$.
- $p_i \leftarrow \text{TFHE.PartDec}(\text{ct}, sk_i)$: Takes as input a ciphertext ct and a secret key sk_i and outputs a partial decryption p_i .

- $\hat{\mu} \leftarrow \text{TFHE.FinDec}(B)$: Takes as input a set $B = \{p_i\}_{i \in S}$ for some $S \subseteq [N]$ and deterministically outputs a message $\hat{\mu} \in \{0, 1, \perp\}$.

Further we remember the definitions of evaluation correctness and simulation security as outlined in, [BEHG20].

Definition 2.2 (Evaluation Correctness [JRS17]). We have that TFHE scheme is correct if for all λ , depth bounds d , access structure \mathbb{A} , circuit $C : \{0, 1\}^k \rightarrow \{0, 1\}$ of depth at most d , $S \in \mathbb{A}$, and $\mu_i \in \{0, 1\}$, we have the following. For $(pk, sk_1, \dots, sk_N) \leftarrow \text{TFHE.Setup}(1^\lambda, 1^d, \mathbb{A})$, $ct_i \leftarrow \text{TFHE.Encrypt}(pk, \mu_i)$ for $i \in [k]$, $\hat{ct} \leftarrow \text{TFHE.Eval}(pk, C, ct_1, \dots, ct_k)$,

$$\Pr [\text{TFHE.FinDec}(pk, \{\text{TFHE.PartDec}(pk, \hat{ct}, sk_i)\}_{i \in S}) = C(\mu_1, \dots, \mu_k)] = 1 - \text{negl}(\lambda).$$

Definition 2.3 (Simulation Security [JRS17]). We say that a TFHE scheme is simulation secure if for all λ , depth bound d , and access structure \mathbb{A} if there exists a stateful PPT simulator, \mathcal{S} , such that for any PPT adversary \mathcal{A} , we have that the experiments $\text{Expt}_{\mathcal{A}, \text{Real}}(1^\lambda, 1^d)$ and $\text{Expt}_{\mathcal{A}, \text{Sim}}(1^\lambda, 1^d)$ are statistically close as a function of λ . The experiments are defined as follows:

- $\text{Expt}_{\mathcal{A}, \text{Real}}(1^\lambda, 1^d)$:
 1. On input the security parameter 1^λ and depth bound d , the adversary outputs $\mathbb{A} \in \mathbb{S}$.
 2. Run $\text{TFHE.Setup}(1^\lambda, 1^d, \mathbb{A})$ to obtain (pk, sk_1, \dots, sk_N) . The adversary is given pk .
 3. The adversary outputs a set $S \subseteq \{P_1, \dots, P_N\}$ such that $S \notin \mathbb{A}$ together with plaintext messages $\mu_1, \dots, \mu_k \in \{0, 1\}$. The adversary is handed over $\{sk_i\}_{i \in S}$
 4. For each μ_i , the adversary is given $\text{TFHE.Encrypt}(pk, \mu_i) \rightarrow ct_i$.
 5. The adversary issues a polynomial number of queries, $(S_i \subseteq \{P_1, \dots, P_N\}, C_i)$. for circuits $C_i : \{0, 1\}^k \rightarrow \{0, 1\}$. After each query the adversary receives for $l \in S_i$ the value

$$\text{TFHE.PartDec}(\text{TFHE.Eval}(C_i, ct_1, \dots, ct_k), sk_l) \rightarrow p_l$$
 6. \mathcal{A} outputs out , the experiment's output.
- $\text{Expt}_{\mathcal{A}, \text{Sim}}(1^\lambda, 1^d)$:
 1. On input the security parameter 1^λ and depth bound d , the adversary outputs $\mathbb{A} \in \mathbb{S}$.
 2. Run $\text{TFHE.Setup}(1^\lambda, 1^d, \mathbb{A})$ to obtain (pk, sk_1, \dots, sk_N) . The adversary is given pk .
 3. \mathcal{A} outputs a set $S^* \subseteq \{P_1, \dots, P_N\}$ such that $S^* \notin \mathbb{A}$ and plaintexts $\mu_1, \dots, \mu_k \in \{0, 1\}$. The simulator is given pk, \mathbb{A}, S^* as input and outputs $\{sk_i\}_{i \in S^*}$ and the state state . The adversary is given $\{sk_i\}_{i \in S^*}$
 4. For each μ_i , the adversary is given $\text{TFHE.Encrypt}(pk, \mu_i) \rightarrow ct_i$.
 5. \mathcal{A} issues a polynomial number of queries of the form $(S_i \subseteq \{P_1, \dots, P_N\}, C_i)$
 6. for circuits $C_i : \{0, 1\}^k \rightarrow \{0, 1\}$. After each query, the simulator computes

$$\text{Sim}_{\text{TFHE}}(C_i, \{ct_l\}_{l=1}^k, C_i(\mu_1, \dots, \mu_k), \text{state}) \rightarrow \{p_l\}_{l \in S_i}$$

and sends $\{p_l\}_{l \in S_i}$ to the adversary.

7. \mathcal{A} outputs out , the experiment's output.

2.2 Data Independent Priority Queue

In this work, we will use data independent queues as studied in [Tof11, MZ14, MDPB23]. Data independent data structures are unique as their control flow and memory access do not depend on input data ([MZ14]).

Definition 2.4 (Word RAM model [MZ14]). In the word RAM model, the RAM has a constant number of public and secret registers and can perform arbitrary operations on a constant number of registers in constant time.

Definition 2.5 (Data Independent Data Structure [MZ14]). In the word RAM model, a data independent data structure is a collection of algorithms where all the algorithms uses RAM such that the RAM can only set its control flow based on registers that are public.

Data independent queues are especially useful as they allow for efficient computation within MPC and FHE as control flow is not dependent on underlying ciphertexts data. We use a data independent queue as outlined in [MDPB23] which allows for

- **PQ.Insert**: Inserts a tag and value, (p, x) into PQ according to the tag's priority.
- **PQ.ExtractFront**: Removes and returns the (p, y) with highest tag priority.
- **PQ.Front**: Returns the (p, y) with highest tag priority without removing the element.

Moreover, we note that the order is stable. I.e. the first inserted among equal tagged elements has a higher priority.

2.3 Reservoir Sampling

Reservoir sampling is an online algorithm which allows for randomly selecting k elements from a stream of n elements while using $\tilde{O}(k)$ space. Algorithm R ([Vit85]) is a simple algorithm which relies on a priority queue with interface:

- **Reservoir.Init**(k) initialize the reservoir sampling algorithm and data structure
- **Reservoir.Insert**($\mu_i, e_i, \text{coin}_i$) where μ_i is i -th item, e_i is independently sampled randomness, and coin_i is a random coin with probability $1/m$ of equaling 1.
 - If $i \leq k$, insert the item into the queue along with e_i as its tag via **PQ.Insert**(e_i, μ_i).
 - If $i > k$ and $\text{coin}_i = 1$, replace the smallest labeled item in the queue with the new item if the coin is 1.
 - If $i > k$ and $\text{coin}_i = 0$, do nothing.
- $\mu_{a_1}, \mu_{a_2}, \dots, \mu_{a_k} \leftarrow \text{Reservoir.Output}()$ where a_1, \dots, a_k are a uniformly random ordered subset of $[n]$
 - Call **PQ.ExtractFront** k times setting μ_{a_ℓ} to the ℓ -th call to **PQ.ExtractFront** where $\ell \in [k]$.

3 TFHE MSLE Protocol

We use a similar notion of ideal functionality for a multi-secret leader election from the ideal functionality of single secret leader election of **LS: CITE**.

The MSLE functionality $\mathcal{F}_{\text{MSLE}}$: Initialize $E, R \leftarrow \emptyset, b \leftarrow 0$. Initialize S to denote the set of sets of active participants in each round. Set $\mathcal{E} \leftarrow \emptyset$ to denote the set of finished elections. Fix some $k \in \mathbb{N}$ to denote the number of rounds. Upon receiving,

- **register** from party P_i , set $R \leftarrow R \cup \{(i, n)\}$, broadcast **(register, i)** to all parties and set $b \leftarrow b + 1$
- **register_elect**(eid, w) from party P_i . If $eid \in \mathcal{E}$ send \perp to P_i and do nothing. If $(i, w) \notin R$ or $(i, w) \in S_{eid}$, send \perp to P_i and do nothing. If S_{eid} is not defined, set $S_{eid} \leftarrow \{(i, w)\}$. Otherwise, set $S_{eid} \leftarrow S_{eid} \cup \{(i, w)\}$.
- **elect**(eid) Elect k leaders from $S_{eid} \subseteq R$ parties. If $|S| \geq k$ and $eid \notin \mathcal{E}$, randomly sample $W^{eid} \subseteq S_{eid}$ where $|W^{eid}| = k$. Then, assign a random ordering to W^{eid} to get ordered set E^{eid} . Next, send **(outcome, eid, a)** to P_j for all $E_a^{eid} = (j, \cdot)$ and **(outcome, eid, \perp)** to P_i if $(i, \cdot) \notin E^{eid}$. Store $E \leftarrow E \cup \{E^{eid}\}$. Set $\mathcal{E} \leftarrow \mathcal{E} \cup \{eid\}$.
- **reveal** from P_i : for $E_{eid} \in E$, if $i = E_{\ell}^{eid}$, broadcast **(result, eid, ℓ, i)**. Otherwise, broadcast **(rejected, eid, ℓ, i)**.

Figure 1: Description of the MSLE functionality, heavily based on the description of SSLE in [LS: CITE Dario and CFG21](#)

3.1 Semi Honest Simulation Security

We can show semi-honest simulation security by showing that the view of each party i in the real protocol can be simulated throughout the course of one election and then that this simulation can be extended for a polynomial number of elections.

Let \bar{n} be the total number of calls to **register_elect**(eid, \cdot, \cdot) before the election for eid , **elect**(eid) is called and that $k \leq \bar{n} \leq n$ as if $\bar{n} < k$, then **elect** returns \perp and is trivial to simulate. Further, let \bar{k} be the total number of calls to **reveal**(eid, \cdot, \cdot) which did not reject, again we have that $\bar{k} \leq k$. We will now show that a simulator can simulate the view of each party i after \bar{n} calls to **register_elect**(eid, \cdot), a call to **elect**(eid), and \bar{k} calls to **reveal**(eid, \cdot) for a fixed eid .

We will now show that Sim_i is indeed a simulator for the view of π_{MSLE} for **register_elect**, **elect**, and **reveal**.

Lemma 3.1. *[register_elect is simulation secure]*

Proof. The simulator gets as input party i 's secret value, s_i , party i 's TFHE share sk_i , eid , and ticket number w . We note that if $eid \in \mathcal{E}$ or $(i, w) \notin R$, then the simulator can simply output \perp and is identical to the real protocol. Otherwise, if party P_i calls **register_elect**, then the simulator knows the input of P_i and can simulate the protocol honestly. If party P_j call **register_elect** for some $j \neq i$, then the view of the protocol is that of $(\text{Enc}_{\text{TFHE}}(c_j), \text{Enc}_{\text{TFHE}}(e_i, \text{coin}_i), \text{Reservoir}^{eid}, \text{Insert}(\text{Enc}_{\text{TFHE}}(c_j, e_i, \text{coin}_i)))$. Note that this view is completely determined by $\text{Enc}_{\text{TFHE}}(c_j)$. Also, note that c_j is drawn from a random distribution and is not in the view of the real protocol. Thus, by semantic security of TFHE (TODO: add citation and refrence to prelis) we have that $\text{Enc}_{\text{TFHE}}(c_j) \stackrel{c}{\equiv} \text{Enc}_{\text{TFHE}}(c'_j)$ where c'_j is a commitment to a random value. \square

Lemma 3.2 (**elect** is simulation secure). *Proof.* If the view of the real protocol is \perp because eid was already called, then the simulator can simply output \perp and is thus identical to the real protocol.

The MSLE Protocol π_{MSLE} : Initialize $b = 0$. Fix some $k \in \mathbb{N}$ to denote the number of rounds. Initialize an empty set of tickets, R . Initialize an empty lookup of sets of parties in each election, S . Initialize an empty lookup of reservoir sampling data structures \mathcal{R} and a set of finished elections \mathcal{E} . Initialize an empty lookup of election results, E .

Initialize $S \leftarrow \emptyset$ to denote the set of sets of active participants in each round.

- **initialize**
 - Set $n = 0$
 - Sample a random TFHE secret key, public key pair sk, pk and publish pk .
 - Sample some secret s_{TFHE} and publish $\text{Enc}_{\text{TFHE}}(s_{\text{TFHE}})$. This will be the key to the PRF
- **register** from party P_i
 - If party P_i does not already have a TFHE share, create a TFHE share, sk_i , for P_i and send the share over a secure channel to P_i .
 - $R \leftarrow R \cup \{(i, b)\}$, set $b \leftarrow b + 1$.
- **register_elect**($eid, \text{Enc}_{\text{TFHE}}(c_i), w$) from party P_i where $c_i = \text{comm}(s_i)$ for a randomly sampled secret s_i .
 - If $eid \in \mathcal{E}$, then send \perp to P_i and do nothing.
 - If $(i, w) \in S_{eid}$ or $(i, w) \notin R$, send \perp and do nothing.
 - Otherwise, if $\mathcal{R}_{eid} \notin \mathcal{R}$, $\mathcal{R}_{eid} \leftarrow \text{Reservoir.Init}(k)$
 - The CRS will be used to run a PRF to get $\text{Enc}_{\text{TFHE}}(e_i), \text{Enc}_{\text{TFHE}}(\text{coin}_i) = \text{Enc}_{\text{TFHE}}(\text{PRF}(c_i, s_{\text{TFHE}}))$ for $i \in S$
 - $\text{Enc}_{\text{TFHE}}(c_i)$ will be fed to the encrypted streaming sampler along with randomness via $\text{Reservoir}^{eid}.\text{Insert}(\text{Enc}_{\text{TFHE}}(c_i), \text{Enc}_{\text{TFHE}}(e_i), \text{Enc}_{\text{TFHE}}(\text{coin}_i))$.
- **elect**(eid)
 - If $eid \in \mathcal{E}$, then return \perp and do nothing.
 - The encrypted streaming sampler will output a list of k messages via calling $\text{Reservoir}^{eid}.\text{Output}()$ k times: $\text{Enc}_{\text{TFHE}}(c_{a_1}), \text{Enc}_{\text{TFHE}}(c_{a_2}), \dots, \text{Enc}_{\text{TFHE}}(c_{a_k})$.
 - Then, at least t parties will submit decryption shares, $p_i = \text{TFHE.PartDec}(\text{Enc}_{\text{TFHE}}(c_{a_1}), \dots, \text{Enc}_{\text{TFHE}}(c_{a_2}))$ to get $c_{a_1}, \dots, c_{a_k} = \text{TFHE.FinDec}(\{p_i\})$.
 - Add $E_{eid} = \{c_{a_1}, \dots, c_{a_k}\}$ to E .
- **reveal**($eid, \ell, \text{Enc}_{\text{TFHE}}(c'_i)$) from P_i
 - P_i submits a proof that they know the opening to c_{a_ℓ} . If this proof verifies, send out $(\text{result}, eid, \ell, i)$ to all parties. Otherwise, send out $(\text{rejected}, eid, \ell, i)$ to all parties.

Figure 2: Description of the MSLE protocol

Simulator for Threshold MSLE Sim_i : Initialize $b = 0$. Fix some $k \in \mathbb{N}$ to denote the number of rounds. Initialize an empty set of tickets, R , an empty lookup of sets of parties in each election, S , a set of finished elections \mathcal{E} , an empty lookup of election results, E , and a set $S \leftarrow \emptyset$ to denote the set of sets of active participants in each round. Moreover, set a random tape for the simulator. The simulator knows, input for party P_i , sk_i and s_i .

- **register** from party P_j
 - $R \leftarrow R \cup \{(i, b)\}$, set $b \leftarrow b + 1$.
 - Simulate the secure channel communication with party P_j if $i \neq j$.
 - Run **register** honestly if $i = j$ by sending sk_i to P_i
- **register_elect**(eid, w, s_i, sk_i) from party P_j .
 - If $j \neq i$, the simulator samples a random value s'_j and follows the protocol directly using s'_{TFHE} and $\text{Enc}_{\text{TFHE}}(c'_j)$ where $c'_j = \text{comm}(s'_j)$. Store c'_j in a lookup table as well as its committed message.
 - If $j = i$, the simulator will use the commitment $c_i = \text{comm}(s_i)$. And run **register_elect** as is in the protocol
- **elect**($eid, \text{out}, s_i, sk_i$) where the **out** is output \perp or $(\text{outcome}, eid, q)_i$ for all $i \in [n]$ where $q \in \{1, \dots, k, \perp\}$.
 - Let \bar{n} be the total number of successful calls to **register_elect**(eid, \dots) before the election for eid
 - If the output is \perp , return \perp .
 - Let a_1, \dots, a_k be the ordered set of parties that won the election. I.e. if $q_i = \ell$ then $a_\ell = i$.
 - The simulator then gives $\text{Sim}_{\text{TFHE}}(C, \{\text{Enc}_{\text{TFHE}}(c'_j)\}_{j \in [\bar{n}]}, c'_1, \dots, c'_{a_\ell}, S, \text{st})$ to the TFHE simulator to get a list of partial decryptions, p_1, \dots, p_{t-1} where S is a qualified set and C is the reservoir sampling circuit with PRF seed s_{TFHE} hardcoded. Note that this sets the decryption of the output of the TFHE circuit sampling to c'_{a_ℓ} for $\ell \in [k]$.
 - Add $E_{eid} = \{c'_{a_1}, \dots, c'_{a_k}\}$ to E .
- **reveal**($eid, \ell, \text{Enc}_{\text{TFHE}}(c'_j)$) from P_j and output \perp or $(\text{result}, eid, \ell, j)$.
 - As the simulator has knowledge of the openings to all of the commitments c'_j used, the simulator to honestly run **reveal**($eid, \ell, \text{Enc}_{\text{TFHE}}(c'_j)$) with commitment c'_j and opening s'_j .
 - Note that in the semi-honest setting, the output is never \perp

Figure 3: Description of the MSLE protocol

Otherwise, note that the simulator needs to “fix” the output of `elect` to yield a list of commitments such that, if $a_\ell = j$ if party j wins election ℓ . Thus, the simulator can simply output the decryption shares $\text{Sim}_{\text{TFHE}}(C, \{\text{Enc}_{\text{TFHE}} c'_j\}_{j \in [\bar{n}]}, c'_1, \dots, c'_{a_\ell}, S, \text{st})$. Note that the reservoir sampling circuit is indeed simulated by Sim_{TFHE} and the decryption shares are simulated as well such that the cipher texts decrypt to $c'_{a_1}, \dots, c'_{a_k} \stackrel{c}{=} c_{a_1}, \dots, c_{a_k}$ by lemma 3.1. \square

Lemma 3.3 (`reveal` is simulation secure). *Proof.* Note that the simulator can simply run `reveal` honestly as the simulator has knowledge of the openings to all of the commitments c'_j used and is thus identical to the view of the real protocol. \square

References

- [BEHG20] Dan Boneh, Saba Eskandarian, Lucjan Hanzlik, and Nicola Greco. Single secret leader election. In *Proceedings of the 2nd ACM Conference on Advances in Financial Technologies*, pages 12–24, 2020. 2.1
- [BGG⁺18] Dan Boneh, Rosario Gennaro, Steven Goldfeder, Aayush Jain, Sam Kim, Peter MR Rasmussen, and Amit Sahai. Threshold cryptosystems from threshold fully homomorphic encryption. In *Advances in Cryptology–CRYPTO 2018: 38th Annual International Cryptology Conference, Santa Barbara, CA, USA, August 19–23, 2018, Proceedings, Part I 38*, pages 565–596. Springer, 2018. 2.1
- [JRS17] Aayush Jain, Peter MR Rasmussen, and Amit Sahai. Threshold fully homomorphic encryption. *Cryptology ePrint Archive*, 2017. 2.1, 2.1, 2.2, 2.3
- [MDPB23] Sahar Mazloom, Benjamin E Diamond, Antigoni Polychroniadou, and Tucker Balch. An efficient data-independent priority queue and its application to dark pools. *Proceedings on Privacy Enhancing Technologies*, 2:5–22, 2023. 2.2, 2.2
- [MZ14] John C Mitchell and Joe Zimmerman. Data-oblivious data structures. In *31st International Symposium on Theoretical Aspects of Computer Science (STACS 2014)*. Schloss Dagstuhl-Leibniz-Zentrum fuer Informatik, 2014. 2.2, 2.4, 2.5
- [Tof11] Tomas Toft. Secure data structures based on multi-party computation. In *Proceedings of the 30th annual ACM SIGACT-SIGOPS symposium on Principles of distributed computing*, pages 291–292, 2011. 2.2
- [Vit85] Jeffrey S Vitter. Random sampling with a reservoir. *ACM Transactions on Mathematical Software (TOMS)*, 11(1):37–57, 1985. 2.3